

ICLR2019

1. STCN: Stochastic Temporal Convolutional Networks
2. Eidetic 3D LSTM: A Model for Video Prediction and Beyond
3. Posterior Attention Models for Sequence to Sequence Learning
4. Bayesian Prediction of Future Street Scenes using Synthetic Likelihoods
5. How Important is a Neuron.
6. Temporal Difference Variational Auto-Encoder
7. Predict then Propagate: Graph Neural Networks meet Personalized PageRank
8. Ordered Neurons: Integrating Tree Structures into Recurrent Neural Networks
9. Pay Less Attention with Lightweight and Dynamic Convolutions
10. Learning protein sequence embeddings using information from structure
11. Learning Implicitly Recurrent CNNs Through Parameter Sharing
12. Capsule Graph Neural Network
13. DyRep: Learning Representations over Dynamic Graphs
14. LanczosNet: Multi-Scale Deep Graph Convolutional Networks
15. Large Scale Graph Learning From Smooth Signals
16. Deep Graph Infomax
17. Supervised Community Detection with Line Graph Neural Networks
18. Invariant and Equivariant Graph Networks
19. On Computation and Generalization of Generative Adversarial Networks under Spectrum Control
20. Adversarial Attacks on Graph Neural Networks via Meta Learning
21. Combining Neural Networks with Personalized PageRank for Classification on Graphs
22. Graph HyperNetworks for Neural Architecture Search
23. Distributional Concavity Regularization for GANs
24. Hyperbolic Attention Networks
25. Structured Neural Summarization
26. PeerNets: Exploiting Peer Wisdom Against Adversarial Attacks
27. Learning deep representations by mutual information estimation and maximization
28. Meta-Learning Update Rules for Unsupervised Representation Learning

29. Learning to Remember More with Less Memorization
30. Rethinking the Value of Network Pruning
31. Learning to Adapt in Dynamic, Real-World Environments through Meta-Reinforcement Learning
32. Conditional Network Embeddings
33. Learning Embeddings into Entropic Wasserstein Spaces
34. Learning Factorized Multimodal Representations
35. What do you learn from context? Probing for sentence structure in contextualized word representations
36. Smoothing the Geometry of Probabilistic Box Embeddings
37. RotatE: Knowledge Graph Embedding by Relational Rotation in Complex Space
38. Measuring Compositionality in Representation Learning
39. Meta-Learning with Latent Embedding Optimization
40. DialogWAE: Multimodal Response Generation with Conditional Wasserstein Auto-Encoder
41. CBOW Is Not All You Need: Combining CBOW with the Compositional Matrix Space Model
42. Understanding Composition of Word Embeddings via Tensor Decomposition
43. Do Deep Generative Models Know What They Don't Know?
44. Understanding and Improving Interpolation in Autoencoders via an Adversarial Regularizer
45. Variational Autoencoders with Jointly Optimized Latent Dependency Structure
46. Poincare Glove: Hyperbolic Word Embeddings
47. Convolutional Neural Networks on Non-uniform Geometrical Signals Using Euclidean Spectral Transformation
48. Diagnosing and Enhancing VAE Models
49. Disjoint Mapping Network for Cross-modal Matching of Voices and Faces
50. Neural Logic Machines
51. RNNs implicitly implement tensor-product representations
52. Generative Code Modeling with Graphs
53. Variational Autoencoders with Jointly Optimized Latent Dependency Structure

54. Spectral Inference Networks: Unifying Deep and Spectral Learning
55. Hindsight policy gradients
56. Wavelet Pooling for Convolutional Neural Networks
57. Deep Convolutional Networks as shallow Gaussian Processes
58. The Deep Weight Prior
59. FlowQA: Grasping Flow in History for Conversational Machine Comprehension
60. On Random Deep Weight-Tied Autoencoders: Exact Asymptotic Analysis, Phase Transitions, and Implications to Trainin
61. GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding
62. textTOvec: DEEP CONTEXTUALIZED NEURAL AUTOREGRESSIVE TOPIC MODELS OF LANGUAGE WITH DISTRIBUTED COMPOSITIONAL PRIOR
63. Don't let your Discriminator be fooled
64. A Kernel Random Matrix-Based Approach for Sparse PCA
65. Identifying and Controlling Important Neurons in Neural Machine Translation
66. Scalable Unbalanced Optimal Transport using Generative Adversarial Networks
67. Multilingual Neural Machine Translation With Soft Decoupled Encoding
68. ProbGAN: Towards Probabilistic GAN with Theoretical Guarantees
69. MARGINALIZED AVERAGE ATTENTIONAL NETWORK FOR WEAKLY-SUPERVISED LEARNING
70. Robust Conditional Generative Adversarial Networks
71. Improving Generalization and Stability of Generative Adversarial Networks
72. Adaptive Input Representations for Neural Language Modeling
73. Supervised Policy Update for Deep Reinforcement Learning
74. Attentive Neural Processes
75. Don't Settle for Average, Go for the Max: Fuzzy Sets and Max-Pooled Word Vectors
76. EvalNE: A Framework for Evaluating Network Embeddings on Link Prediction

77. Correlated Variational Auto-Encoders
78. Context Mover's Distance & Barycenters Optimal transport of contexts for building representations
79. Learning Embeddings into Entropic Wasserstein Spaces
80. Universal Deep Node Representation via Partial Permutation Invariant Set Functions
81. Graph U-Net
82. The Laplacian in RL: Learning Representations with Efficient Approximations
83. BIGSAGE: unsupervised inductive representation learning of graph via bi-attended sampling and global-biased aggregating
84. DEEP GEOMETRICAL GRAPH CLASSIFICATION
85. CGNF: Conditional Graph Neural Fields
86. Link Prediction in Hypergraphs using Graph Convolutional Networks
87. Learning Graph Decomposition
88. Supervised Contextual Embeddings for Transfer Learning in Natural Language Processing Tasks
89. Convolutional Neural Networks on Non-uniform Geometrical Signals Using Euclidean Spectral Transformation
90. On Self Modulation for Generative Adversarial Networks

91. RelGAN: Relational Generative Adversarial Networks for Text Generation
92. Generating High Fidelity Images with Subscale Pixel Networks and Multidimensional Upscaling
93. Sliced Wasserstein Auto-Encoders
94. Improving MMD-GAN Training with Repulsive Loss Function
95. The relativistic discriminator: a key element missing from standard GAN
96. Diversity-Sensitive Conditional Generative Adversarial Networks
97. Unsupervised Adversarial Image Reconstruction
98. Discriminator Rejection Sampling
99. Towards GAN Benchmarks Which Require Generalization
100. GAN Dissection: Visualizing and Understanding Generative Adversarial Networks