

- [1] MixHop: Higher-Order Graph Convolutional Architectures via Sparsified Neighborhood Mixing
- [2] Multi-objective training of Generative Adversarial Networks with multiple discriminators
- [3] Graph Element Networks: adaptive, structured computation and memory
- [4] Analogies Explained: Towards Understanding Word Embeddings
- [5] Fine-Grained Analysis of Optimization and Generalization for Overparameterized Two-Layer Neural Networks
- [6] Entropic GANs meet VAEs: A Statistical Approach to Compute Sample Likelihoods in GANs
- [7] Provable Guarantees for Gradient-Based Meta-Learning
- [8] Concrete Autoencoders: Differentiable Feature Selection and Reconstruction
- [9] Learning to Route in Similarity Graphs
- [10] Active Learning for Probabilistic Structured Prediction of Cuts and Matchings
- [11] Invertible Residual Networks
- [12] Overcoming Multi-model Forgetting
- [13] Optimal Kronecker-Sum Approximation of Real Time Recurrent Learning
- [14] A Kernel Perspective for Regularizing Deep Neural Networks
- [15] Adversarial Attacks on Node Embeddings via Graph Poisoning
- [16] Compositional Fairness Constraints for Graph Embeddings
- [17] Understanding the Origins of Bias in Word Embeddings
- [18] Why do Larger Models Generalize Better? A Theoretical Perspective via the XOR Problem
- [19] What is the Effect of Importance Weighting in Deep Learning?
- [20] Learning Action Representations for Reinforcement Learning
- [21] Dynamic Measurement Scheduling for Event Forecasting using Deep RL
- [22] Neural Network Attributions: A Causal Perspective
- [23] Information-Theoretic Considerations in Batch Reinforcement Learning
- [24] Generative Adversarial User Model for Reinforcement Learning Based Recommendation System
- [25] Understanding and Utilizing Deep Neural Networks Trained with Noisy Labels
- [26] A Gradual, Semi-Discrete Approach to Generative Network Training via Explicit Wasserstein Minimization
- [27] Transferability vs. Discriminability: Batch Spectral Penalization for Adversarial Domain Adaptation
- [28] Fast Incremental von Neumann Graph Entropy Computation: Theory, Algorithm, and Applications
- [29] Multivariate-Information Adversarial Ensemble for Scalable Joint Distribution Matching
- [30] Robust Decision Trees Against Adversarial Examples
- [31] RaFM: Rank-Aware Factorization Machines
- [32] Control Regularization for Reduced Variance Reinforcement Learning
- [33] Random Walks on Hypergraphs with Edge-Dependent Vertex Weights
- [34] Probability Functional Descent: A Unifying Perspective on GANs, Variational Inference, and Reinforcement Learning
- [35] MeanSum: A Neural Model for Unsupervised Multi-Document Abstractive

## Summarization

- [36] Sensitivity Analysis of Linear Structural Causal Models
- [37] Quantifying Generalization in Reinforcement Learning
- [38] Empirical Analysis of Beam Search Performance Degradation in Neural Sequence Models
- [39] Certified Adversarial Robustness via Randomized Smoothing
- [40] Gauge Equivariant Convolutional Networks and the Icosahedral CNN
- [41] CURIOUS: Intrinsically Motivated Modular Multi-Goal Reinforcement Learning
- [42] A fully differentiable beam search decoder
- [43] Monge blunts Bayes: Hardness Results for Adversarial Training
- [44] Flexibly Fair Representation Learning by Disentanglement
- [45] Teaching a black-box learner
- [46] Stochastic Deep Networks
- [47] Learning to Convolve: A Generalized Weight-Tying Approach
- [48] Generalized No Free Lunch Theorem for Adversarial Robustness
- [49] Wasserstein of Wasserstein Loss for Learning Generative Models
- [50] GDPP: Learning Diverse Generations using Determinantal Point Processes
- [51] On the Connection Between Adversarial Robustness and Saliency Map Interpretability
- [52] Dead-ends and Secure Exploration in Reinforcement Learning
- [53] DL2: Training and Querying Neural Networks with Logic
- [54] On discriminative learning of prediction uncertainty
- [55] Learning Discrete Structures for Graph Neural Networks
- [56] Transfer Learning for Related Reinforcement Learning Tasks via Image-to-Image Translation
- [57] Graph U-Nets
- [58] Deep Generative Learning via Variational Gradient Flow
- [59] Geometric Scattering for Graph Data Analysis
- [60] Multi-Frequency Phase Synchronization
- [61] SelectiveNet: A Deep Neural Network with an Integrated Reject Option
- [62] Recursive Sketches for Modular Deep Learning
- [63] Adversarial Examples Are a Natural Consequence of Test Error in Noise
- [64] Estimating Information Flow in Deep Neural Networks
- [65] Efficient Training of BERT by Progressively Stacking
- [66] A Statistical Investigation of Long Memory in Language and Music
- [67] Graphite: Iterative Generative Modeling of Graphs
- [68] Fast Algorithm for Generalized Multinomial Models with Ranking Data
- [69] Towards a Deep and Unified Understanding of Deep Neural Models in NLP
- [70] Humor in Word Embeddings: Cockamamie Gobbledegook for Nincompoops
- [71] Exploring interpretable LSTM neural networks over multi-variable data
- [72] Learning to Exploit Long-term Relational Dependencies in Knowledge Graphs
- [73] IMEXnet A Forward Stable Deep Neural Network
- [74] Understanding and Controlling Memory in Recurrent Neural Networks
- [75] On the Impact of the Activation function on Deep Neural Networks Training
- [76] Using Pre-Training Can Improve Model Robustness and Uncertainty

- [77] Parameter-Efficient Transfer Learning for NLP
- [78] Finding Mixed Nash Equilibria of Generative Adversarial Networks
- [79] Hierarchical Importance Weighted Autoencoders
- [80] HexaGAN: Generative Adversarial Nets for Real World Classification
- [81] Learning What and Where to Transfer
- [82] Graph Neural Network for Music Score Data and Modeling Expressive Piano Performance
- [83] Ladder Capsule Network
- [84] Learning Discrete and Continuous Factors of Data via Alternating Disentanglement
- [85] Neural Logic Reinforcement Learning
- [86] Kernel Mean Matching for Content Addressability of GANs
- [87] Trainable Decoding of Sets of Sequences for Neural Sequence Models
- [88] POPQORN: Quantifying Robustness of Recurrent Neural Networks
- [89] Stochastic Beams and Where To Find Them: The Gumbel-Top-k Trick for Sampling Sequences Without Replacement
- [90] Similarity of Neural Network Representations Revisited
- [91] Characterizing Well-Behaved vs. Pathological Deep Neural Networks
- [92] State-Reification Networks: Improving Generalization by Modeling the Distribution of Hidden Representations
- [93] Self-Attention Graph Pooling
- [94] Set Transformer: A Framework for Attention-based Permutation-Invariant Neural Networks
- [95] Are Generative Classifiers More Robust to Adversarial Attacks?
- [96] LGM-Net: Learning to Generate Matching Networks for Few-Shot Learning
- [97] Graph Matching Networks for Learning the Similarity of Graph Structured Objects
- [98] Learn to Grow: A Continual Structure Learning Framework for Overcoming Catastrophic Forgetting
- [99] Challenging Common Assumptions in the Unsupervised Learning of Disentangled Representations
- [100] Disentangled Graph Convolutional Networks
- [101] Traditional and Heavy Tailed Self Regularization in Neural Network Models
- [102] Disentangling Disentanglement in Variational Autoencoders
- [103] MIWAE: Deep Generative Modelling and Imputation of Incomplete Data Sets
- [104] Imputing Missing Events in Continuous-Time Event Streams
- [105] The Wasserstein Transform
- [106] Relational Pooling for Graph Representations
- [107] Dropout as a Structured Shrinkage Prior
- [108] Zero-Shot Knowledge Distillation in Deep Networks
- [109] Improving Adversarial Robustness via Promoting Ensemble Diversity
- [110] Deep Residual Output Layers for Neural Language Generation
- [111] Domain Agnostic Learning with Disentangled Representations
- [112] Towards Understanding Knowledge Distillation
- [113] GMNN: Graph Markov Neural Networks
- [114] HyperGAN: A Generative Model for Diverse, Performant Neural Networks

- [115] Mixture Models for Diverse Machine Translation: Tricks of the Trade
- [116] Revisiting precision recall definition for generative modeling
- [117] The Evolved Transformer
- [118] MASS: Masked Sequence to Sequence Pre-training for Language Generation
- [119] Insertion Transformer: Flexible Sequence Generation via Insertion Operations
- [120] BERT and PALs: Projected Attention Layers for Efficient Adaptation in Multi-Task Learning
- [121] Contextual Memory Trees
- [122] Robustly Disentangled Causal Mechanisms: Validating Deep Representations for Interventional Robustness
- [123] Hyperbolic Disk Embeddings for Directed Acyclic Graphs
- [124] Equivariant Transformer Networks
- [125] Making Deep Q-learning methods robust to time discretization
- [126] The Natural Language of Actions
- [127] Optimal Transport for structured data with application on graphs
- [128] Probabilistic Neural Symbolic Models for Interpretable Visual Question Answering
- [129] SATNet: Bridging deep learning and logical reasoning using a differentiable satisfiability solver
- [130] Improving Neural Language Modeling via Adversarial Training
- [131] State-Regularized Recurrent Neural Networks
- [132] Bias Also Matters: Bias Attribution for Deep Neural Network Explanation
- [133] Jumpout : Improved Dropout for Deep Neural Networks with ReLUs
- [134] Non-Monotonic Sequential Text Generation
- [135] Simplifying Graph Convolutional Networks
- [136] Gromov-Wasserstein Learning for Graph Matching and Node Embedding
- [137] Position-aware Graph Neural Networks
- [138] Learning Neurosymbolic Generative Models via Program Synthesis
- [139] DAG-GNN: DAG Structure Learning with Graph Neural Networks
- [140] Self-Attention Generative Adversarial Networks